

tags: **dlcv**

HW2

P1

```
StyleGAN2(
    (S): StyleVectorizer(
        (net): Sequential(
            (0): EqualLinear()
            (1): LeakyReLU(negative_slope=0.2, inplace=True)
            (2): EqualLinear()
            (3): LeakyReLU(negative_slope=0.2, inplace=True)
            (4): EqualLinear()
            (5): LeakyReLU(negative_slope=0.2, inplace=True)
            (6): EqualLinear()
            (7): LeakyReLU(negative_slope=0.2, inplace=True)
            (8): EqualLinear()
            (9): LeakyReLU(negative_slope=0.2, inplace=True)
            (10): EqualLinear()
            (11): LeakyReLU(negative_slope=0.2, inplace=True)
            (12): EqualLinear()
            (13): LeakyReLU(negative_slope=0.2, inplace=True)
            (14): EqualLinear()
            (15): LeakyReLU(negative_slope=0.2, inplace=True)
        )
    )
    (G): Generator(
        (initial_conv): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding
        (blocks): ModuleList(
            (0): GeneratorBlock(
                (to_style1): Linear(in_features=512, out_features=512, bias=True)
                (to_noise1): Linear(in_features=1, out_features=512, bias=True)
                (conv1): Conv2DMod()
                (to_style2): Linear(in_features=512, out_features=512, bias=True)
                (to_noise2): Linear(in_features=1, out_features=512, bias=True)
                (conv2): Conv2DMod()
                (activation): LeakyReLU(negative_slope=0.2, inplace=True)
                (to_rgb): RGBBlock(
                    (to_style): Linear(in_features=512, out_features=512, bias=True)
                    (conv): Conv2DMod()
                    (upsample): Sequential(
                        (0): Upsample(scale_factor=2.0, mode=bilinear)
                        (1): Blur()
                    )
                )
            )
        )
        (1): GeneratorBlock(
            (upsample): Upsample(scale_factor=2.0, mode=bilinear)
            (to_style1): Linear(in_features=512, out_features=512, bias=True)
            (to_noise1): Linear(in_features=1, out_features=256, bias=True)
            (conv1): Conv2DMod()
            (to_style2): Linear(in_features=512, out_features=256, bias=True)
            (to_noise2): Linear(in_features=1, out_features=256, bias=True)
            (conv2): Conv2DMod()
            (activation): LeakyReLU(negative_slope=0.2, inplace=True)
            (to_rgb): RGBBlock(
                (to_style): Linear(in_features=512, out_features=256, bias=True)

```

```
(conv): Conv2DMod()
(upsample): Sequential(
    (0): Upsample(scale_factor=2.0, mode=bilinear)
    (1): Blur()
)
)
(2): GeneratorBlock(
    (upsample): Upsample(scale_factor=2.0, mode=bilinear)
    (to_style1): Linear(in_features=512, out_features=256, bias=True)
    (to_noise1): Linear(in_features=1, out_features=128, bias=True)
    (conv1): Conv2DMod()
    (to_style2): Linear(in_features=512, out_features=128, bias=True)
    (to_noise2): Linear(in_features=1, out_features=128, bias=True)
    (conv2): Conv2DMod()
    (activation): LeakyReLU(negative_slope=0.2, inplace=True)
    (to_rgb): RGBBlock(
        (to_style): Linear(in_features=512, out_features=128, bias=True)
        (conv): Conv2DMod()
        (upsample): Sequential(
            (0): Upsample(scale_factor=2.0, mode=bilinear)
            (1): Blur()
        )
    )
)
(3): GeneratorBlock(
    (upsample): Upsample(scale_factor=2.0, mode=bilinear)
    (to_style1): Linear(in_features=512, out_features=128, bias=True)
    (to_noise1): Linear(in_features=1, out_features=64, bias=True)
    (conv1): Conv2DMod()
    (to_style2): Linear(in_features=512, out_features=64, bias=True)
    (to_noise2): Linear(in_features=1, out_features=64, bias=True)
    (conv2): Conv2DMod()
    (activation): LeakyReLU(negative_slope=0.2, inplace=True)
    (to_rgb): RGBBlock(
        (to_style): Linear(in_features=512, out_features=64, bias=True)
        (conv): Conv2DMod()
        (upsample): Sequential(
            (0): Upsample(scale_factor=2.0, mode=bilinear)
            (1): Blur()
        )
    )
)
(4): GeneratorBlock(
    (upsample): Upsample(scale_factor=2.0, mode=bilinear)
    (to_style1): Linear(in_features=512, out_features=64, bias=True)
    (to_noise1): Linear(in_features=1, out_features=32, bias=True)
    (conv1): Conv2DMod()
    (to_style2): Linear(in_features=512, out_features=32, bias=True)
    (to_noise2): Linear(in_features=1, out_features=32, bias=True)
    (conv2): Conv2DMod()
    (activation): LeakyReLU(negative_slope=0.2, inplace=True)
    (to_rgb): RGBBlock()
```

```
(to_style): Linear(in_features=512, out_features=32, bias=True)
(conv): Conv2DMod()
)
)
)
(attns): ModuleList(
(0): None
(1): None
(2): None
(3): Sequential(
(0): Residual(
(fn): PreNorm(
(fn): LinearAttention(
(nonlin): GELU()
(to_q): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
(to_kv): DepthWiseConv2d(
(net): Sequential(
(0): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=1)
(1): Conv2d(128, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
)
)
(to_out): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1))
)
(norm): ChanNorm()
)
)
(1): Residual(
(fn): PreNorm(
(fn): Sequential(
(0): Conv2d(128, 256, kernel_size=(1, 1), stride=(1, 1))
(1): LeakyReLU(negative_slope=0.2, inplace=True)
(2): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1))
)
(norm): ChanNorm()
)
)
)
)
(4): Sequential(
(0): Residual(
(fn): PreNorm(
(fn): LinearAttention(
(nonlin): GELU()
(to_q): Conv2d(64, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
(to_kv): DepthWiseConv2d(
(net): Sequential(
(0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=1)
(1): Conv2d(64, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
)
)
(to_out): Conv2d(512, 64, kernel_size=(1, 1), stride=(1, 1))
)
(norm): ChanNorm()
)
)
```

```
)  
    (1): Residual(  
        (fn): PreNorm(  
            (fn): Sequential(  
                (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(1, 1))  
                (1): LeakyReLU(negative_slope=0.2, inplace=True)  
                (2): Conv2d(128, 64, kernel_size=(1, 1), stride=(1, 1))  
            )  
            (norm): ChanNorm()  
        )  
    )  
)  
(SE): StyleVectorizer(  
    (net): Sequential(  
        (0): EqualLinear()  
        (1): LeakyReLU(negative_slope=0.2, inplace=True)  
        (2): EqualLinear()  
        (3): LeakyReLU(negative_slope=0.2, inplace=True)  
        (4): EqualLinear()  
        (5): LeakyReLU(negative_slope=0.2, inplace=True)  
        (6): EqualLinear()  
        (7): LeakyReLU(negative_slope=0.2, inplace=True)  
        (8): EqualLinear()  
        (9): LeakyReLU(negative_slope=0.2, inplace=True)  
        (10): EqualLinear()  
        (11): LeakyReLU(negative_slope=0.2, inplace=True)  
        (12): EqualLinear()  
        (13): LeakyReLU(negative_slope=0.2, inplace=True)  
        (14): EqualLinear()  
        (15): LeakyReLU(negative_slope=0.2, inplace=True)  
    )  
)  
(GE): Generator(  
    (initial_conv): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding  
    (blocks): ModuleList(  
        (0): GeneratorBlock(  
            (to_style1): Linear(in_features=512, out_features=512, bias=True)  
            (to_noise1): Linear(in_features=1, out_features=512, bias=True)  
            (conv1): Conv2DMod()  
            (to_style2): Linear(in_features=512, out_features=512, bias=True)  
            (to_noise2): Linear(in_features=1, out_features=512, bias=True)  
            (conv2): Conv2DMod()  
            (activation): LeakyReLU(negative_slope=0.2, inplace=True)  
            (to_rgb): RGBBlock(  
                (to_style): Linear(in_features=512, out_features=512, bias=True)  
                (conv): Conv2DMod()  
                (upsample): Sequential(  
                    (0): Upsample(scale_factor=2.0, mode=bilinear)  
                    (1): Blur()  
                )  
            )  
        )  
    )
```

```
)  
(1): GeneratorBlock(  
    (upsample): Upsample(scale_factor=2.0, mode=bilinear)  
    (to_style1): Linear(in_features=512, out_features=512, bias=True)  
    (to_noise1): Linear(in_features=1, out_features=256, bias=True)  
    (conv1): Conv2DMod()  
    (to_style2): Linear(in_features=512, out_features=256, bias=True)  
    (to_noise2): Linear(in_features=1, out_features=256, bias=True)  
    (conv2): Conv2DMod()  
    (activation): LeakyReLU(negative_slope=0.2, inplace=True)  
    (to_rgb): RGBBlock(  
        (to_style): Linear(in_features=512, out_features=256, bias=True)  
        (conv): Conv2DMod()  
        (upsample): Sequential(  
            (0): Upsample(scale_factor=2.0, mode=bilinear)  
            (1): Blur()  
        )  
    )  
)  
(2): GeneratorBlock(  
    (upsample): Upsample(scale_factor=2.0, mode=bilinear)  
    (to_style1): Linear(in_features=512, out_features=256, bias=True)  
    (to_noise1): Linear(in_features=1, out_features=128, bias=True)  
    (conv1): Conv2DMod()  
    (to_style2): Linear(in_features=512, out_features=128, bias=True)  
    (to_noise2): Linear(in_features=1, out_features=128, bias=True)  
    (conv2): Conv2DMod()  
    (activation): LeakyReLU(negative_slope=0.2, inplace=True)  
    (to_rgb): RGBBlock(  
        (to_style): Linear(in_features=512, out_features=128, bias=True)  
        (conv): Conv2DMod()  
        (upsample): Sequential(  
            (0): Upsample(scale_factor=2.0, mode=bilinear)  
            (1): Blur()  
        )  
    )  
)  
(3): GeneratorBlock(  
    (upsample): Upsample(scale_factor=2.0, mode=bilinear)  
    (to_style1): Linear(in_features=512, out_features=128, bias=True)  
    (to_noise1): Linear(in_features=1, out_features=64, bias=True)  
    (conv1): Conv2DMod()  
    (to_style2): Linear(in_features=512, out_features=64, bias=True)  
    (to_noise2): Linear(in_features=1, out_features=64, bias=True)  
    (conv2): Conv2DMod()  
    (activation): LeakyReLU(negative_slope=0.2, inplace=True)  
    (to_rgb): RGBBlock(  
        (to_style): Linear(in_features=512, out_features=64, bias=True)  
        (conv): Conv2DMod()  
        (upsample): Sequential(  
            (0): Upsample(scale_factor=2.0, mode=bilinear)  
            (1): Blur()  
        )
```

```
)  
)  
(4): GeneratorBlock(  
    (upsample): Upsample(scale_factor=2.0, mode=bilinear)  
    (to_style1): Linear(in_features=512, out_features=64, bias=True)  
    (to_noise1): Linear(in_features=1, out_features=32, bias=True)  
    (conv1): Conv2DMod()  
    (to_style2): Linear(in_features=512, out_features=32, bias=True)  
    (to_noise2): Linear(in_features=1, out_features=32, bias=True)  
    (conv2): Conv2DMod()  
    (activation): LeakyReLU(negative_slope=0.2, inplace=True)  
    (to_rgb): RGBBlock(  
        (to_style): Linear(in_features=512, out_features=32, bias=True)  
        (conv): Conv2DMod()  
    )  
)  
(attns): ModuleList(  
    (0): None  
    (1): None  
    (2): None  
    (3): Sequential(  
        (0): Residual(  
            (fn): PreNorm(  
                (fn): LinearAttention(  
                    (nonlin): GELU()  
                    (to_q): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=  
                    (to_kv): DepthWiseConv2d(  
                        (net): Sequential(  
                            (0): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padd  
                            (1): Conv2d(128, 1024, kernel_size=(1, 1), stride=(1, 1), bia  
                        )  
                    )  
                    (to_out): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1))  
                )  
                (norm): ChanNorm()  
            )  
        )  
        (1): Residual(  
            (fn): PreNorm(  
                (fn): Sequential(  
                    (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(1, 1))  
                    (1): LeakyReLU(negative_slope=0.2, inplace=True)  
                    (2): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1))  
                )  
                (norm): ChanNorm()  
            )  
        )  
    )  
(4): Sequential(  
    (0): Residual(  
        (fn): PreNorm(  
            (fn): LinearAttention(  
)
```

```
(nonlin): GELU()
(to_q): Conv2d(64, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
(to_kv): DepthWiseConv2d(
    (net): Sequential(
        (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=1)
        (1): Conv2d(64, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    )
    (to_out): Conv2d(512, 64, kernel_size=(1, 1), stride=(1, 1))
)
(norm): ChanNorm()
)
)
(1): Residual(
    (fn): PreNorm(
        (fn): Sequential(
            (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(1, 1))
            (1): LeakyReLU(negative_slope=0.2, inplace=True)
            (2): Conv2d(128, 64, kernel_size=(1, 1), stride=(1, 1))
        )
        (norm): ChanNorm()
    )
)
)
)
)
```

I used stylegan2 with 2 attention layers, and trained 150000 steps. However the checkpoints around 8200~9200 steps generate best images.

2.



3~4.

IS:2.1766

FID:27.3224

5.

I tried DCGAN, WGAN-GP and SNGAN. SNGAN > WGAN-GP > DCGAN, but stylegan is the best.

In stylegan, when generating images, the trunc_psi should be closer to 1.0 to have better image diversity.

P2

1.

I built a generator and discriminator where both of them have convolutional layers. And my discriminator will return 2 results, label results and validity results. The image should be resized to 32, otherwise the accuracy will be low.

```
Generator(  
    (label_emb): Embedding(10, 100)  
    (l1): Sequential(  
        (0): Linear(in_features=100, out_features=8192, bias=True)  
    )  
    (conv_blocks): Sequential(  
        (0): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_s  
        (1): Upsample(scale_factor=2.0, mode=nearest)  
  
        (2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (3): BatchNorm2d(128, eps=0.8, momentum=0.1, affine=True, track_running_st  
        (4): LeakyReLU(negative_slope=0.2, inplace=True)  
        (5): Upsample(scale_factor=2.0, mode=nearest)  
        (6): Conv2d(128, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (7): BatchNorm2d(64, eps=0.8, momentum=0.1, affine=True, track_running_stat  
        (8): LeakyReLU(negative_slope=0.2, inplace=True)  
        (9): Conv2d(64, 3, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (10): Tanh()  
    )  
)
```

```
Discriminator(  
    (conv_blocks): Sequential(  
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))  
        (1): LeakyReLU(negative_slope=0.2, inplace=True)  
        (2): Dropout2d(p=0.25, inplace=False)  
        (3): Conv2d(16, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))  
        (4): LeakyReLU(negative_slope=0.2, inplace=True)  
        (5): Dropout2d(p=0.25, inplace=False)  
        (6): BatchNorm2d(32, eps=0.8, momentum=0.1, affine=True, track_running_stat  
        (7): Conv2d(32, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))  
        (8): LeakyReLU(negative_slope=0.2, inplace=True)  
        (9): Dropout2d(p=0.25, inplace=False)  
        (10): BatchNorm2d(64, eps=0.8, momentum=0.1, affine=True, track_running_st  
        (11): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))  
        (12): LeakyReLU(negative_slope=0.2, inplace=True)  
        (13): Dropout2d(p=0.25, inplace=False)  
        (14): BatchNorm2d(128, eps=0.8, momentum=0.1, affine=True, track_running_st  
    )  
    (adv_layer): Sequential(  
        (0): Linear(in_features=512, out_features=1, bias=True)  
        (1): Sigmoid()  
    )  
    (aux_layer): Sequential(  
        (0): Linear(in_features=512, out_features=10, bias=True)  
        (1): Softmax(dim=0)  
    )  
)
```

2~4.

0.95

5.



P3

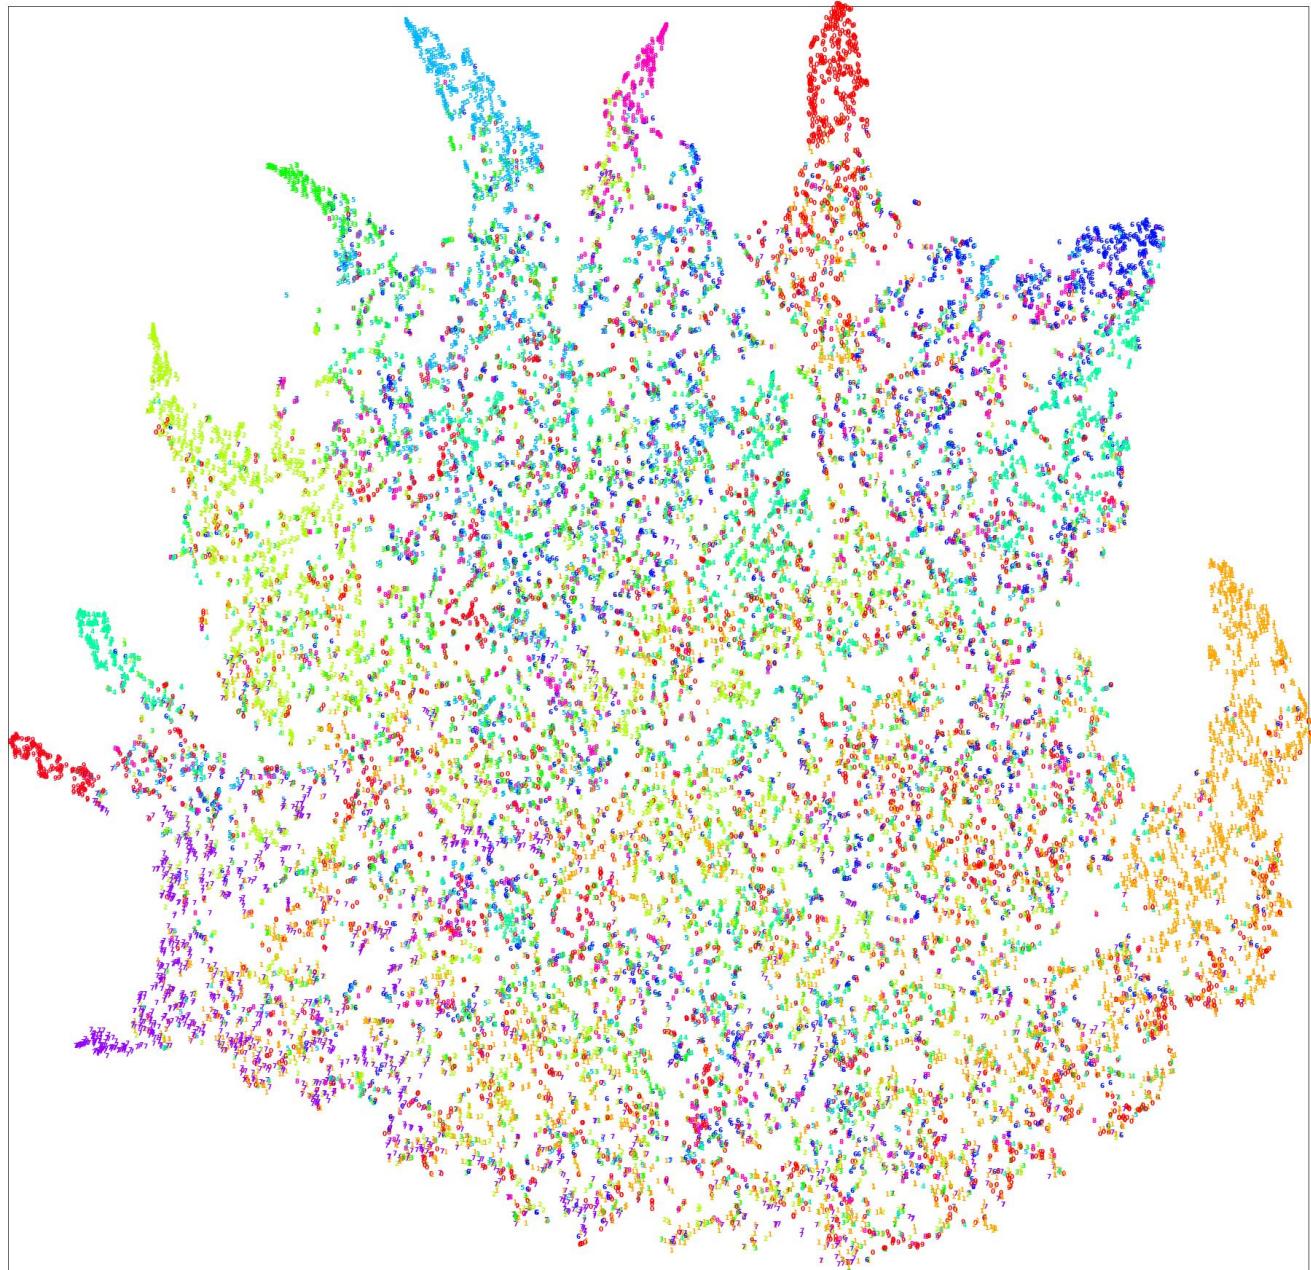
1~3 .

	MNSITM → USPS	SVHN → MNIST-M	USPS → SVHN
On source	0.7045	0.4443	0.2127
DANN	0.7643	0.4918	0.3041
target	0.9583	0.9513	0.8722

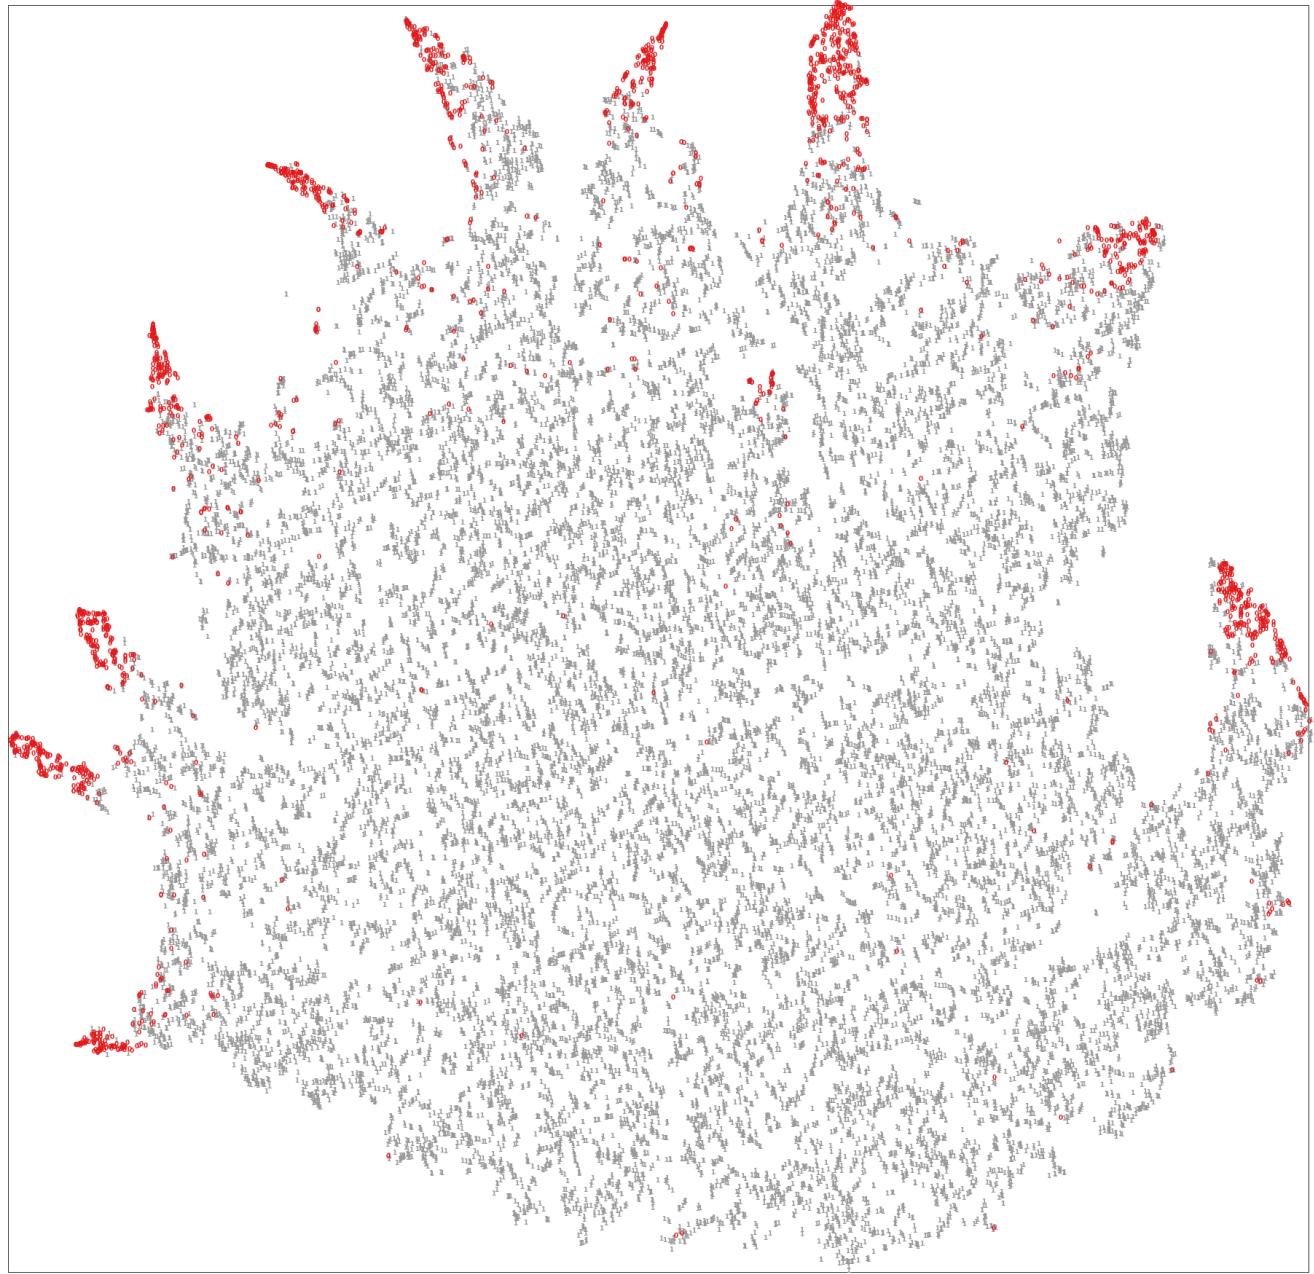
4.

- usps->svhn

(a)

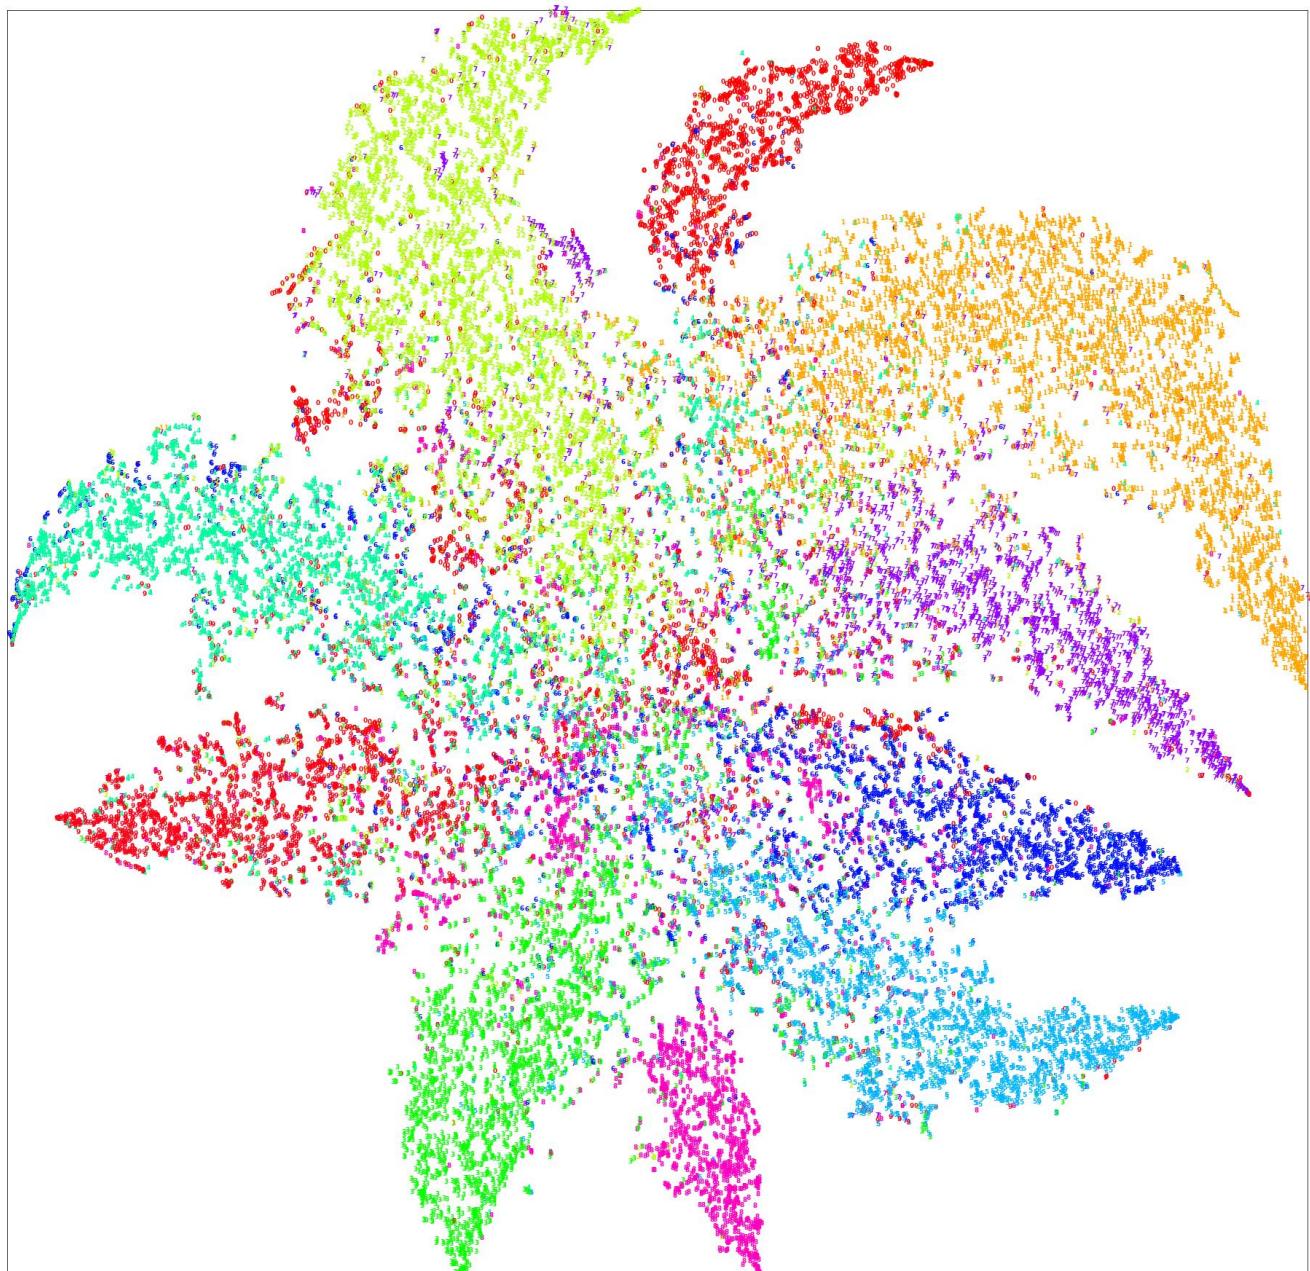


(b)

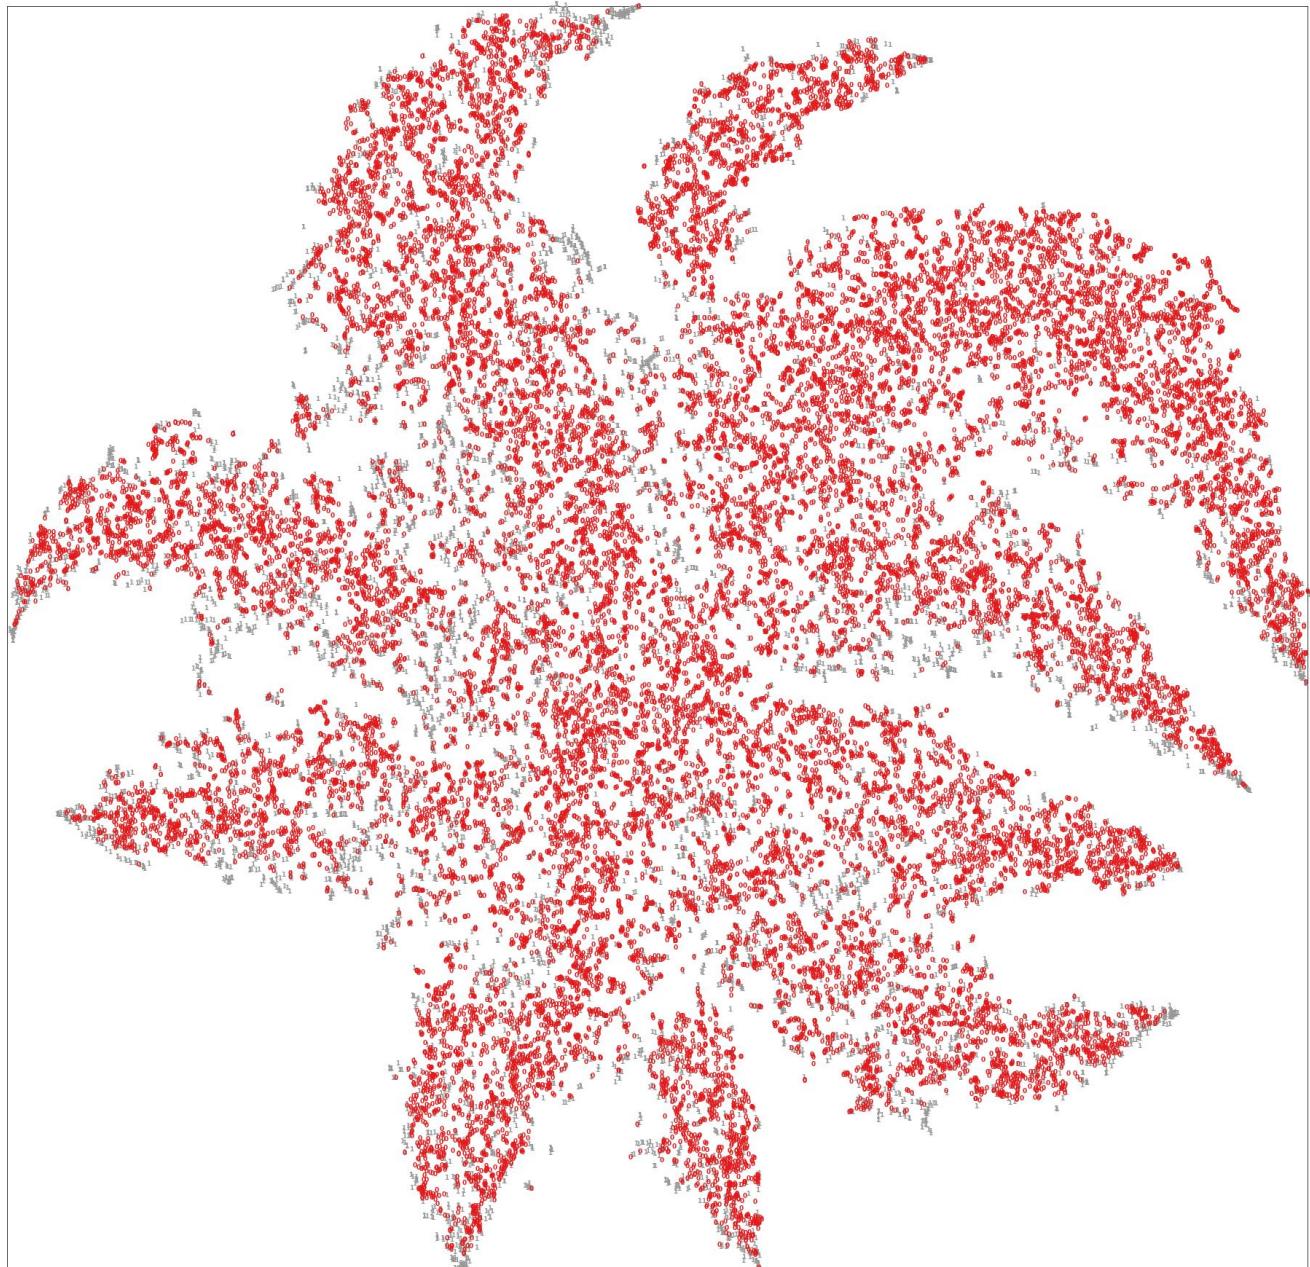


- svhn->mnistm

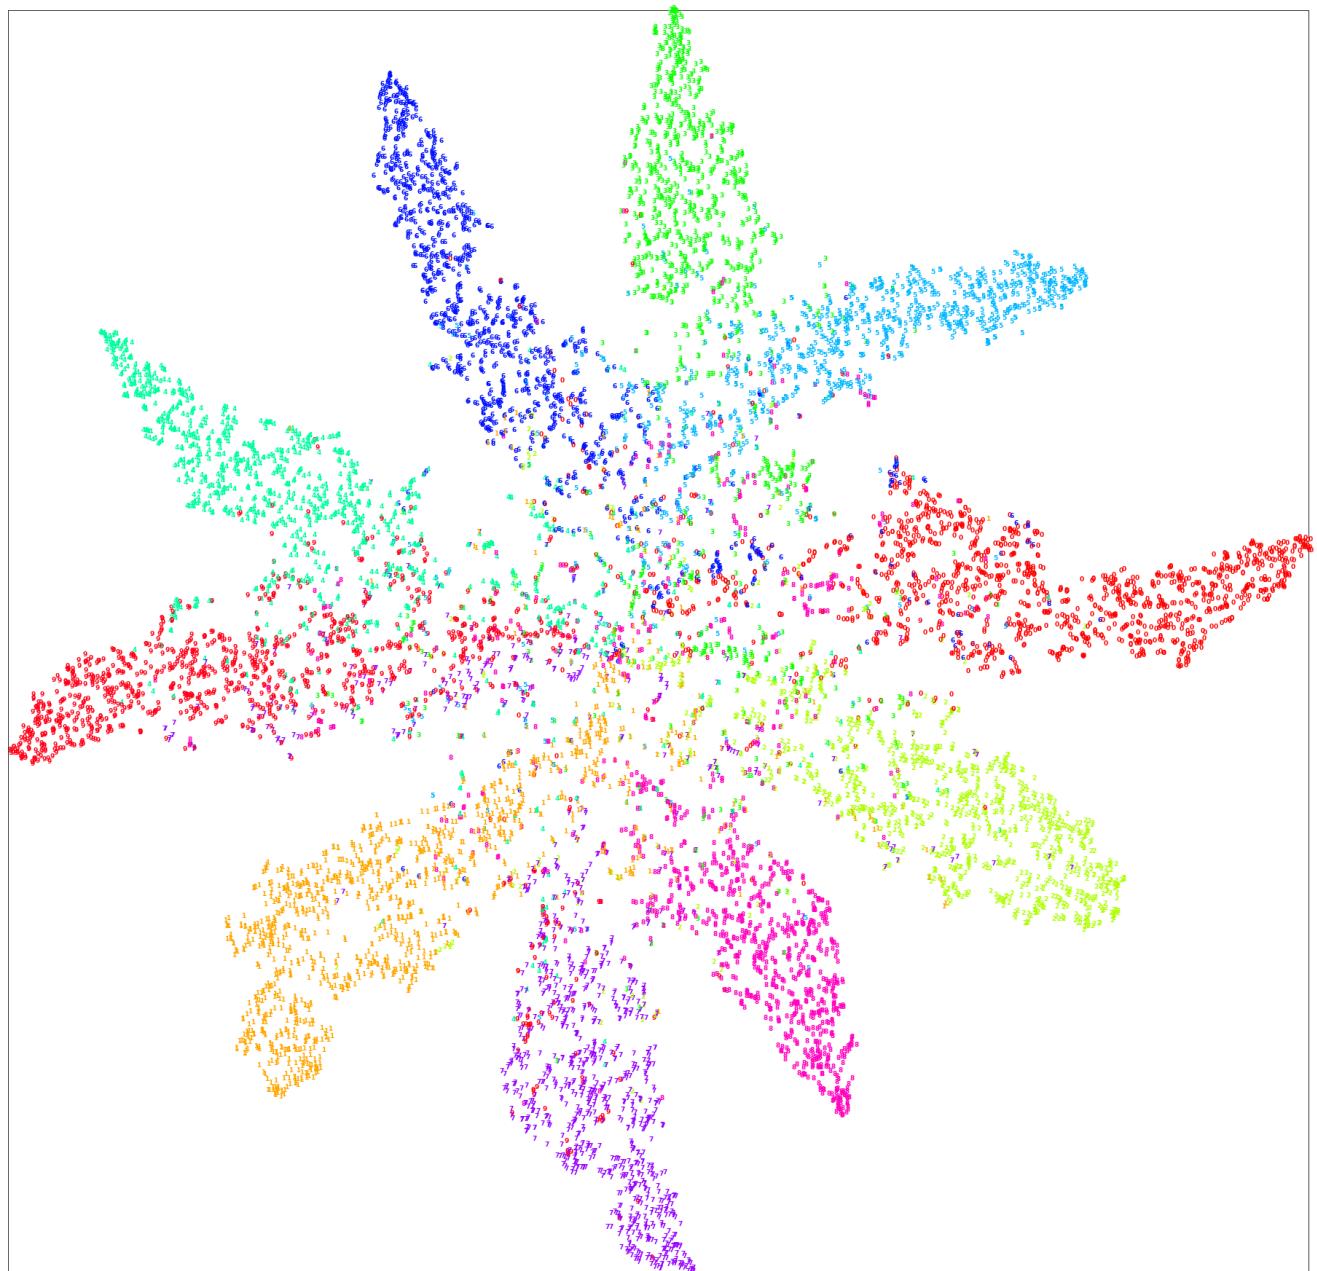
(a)



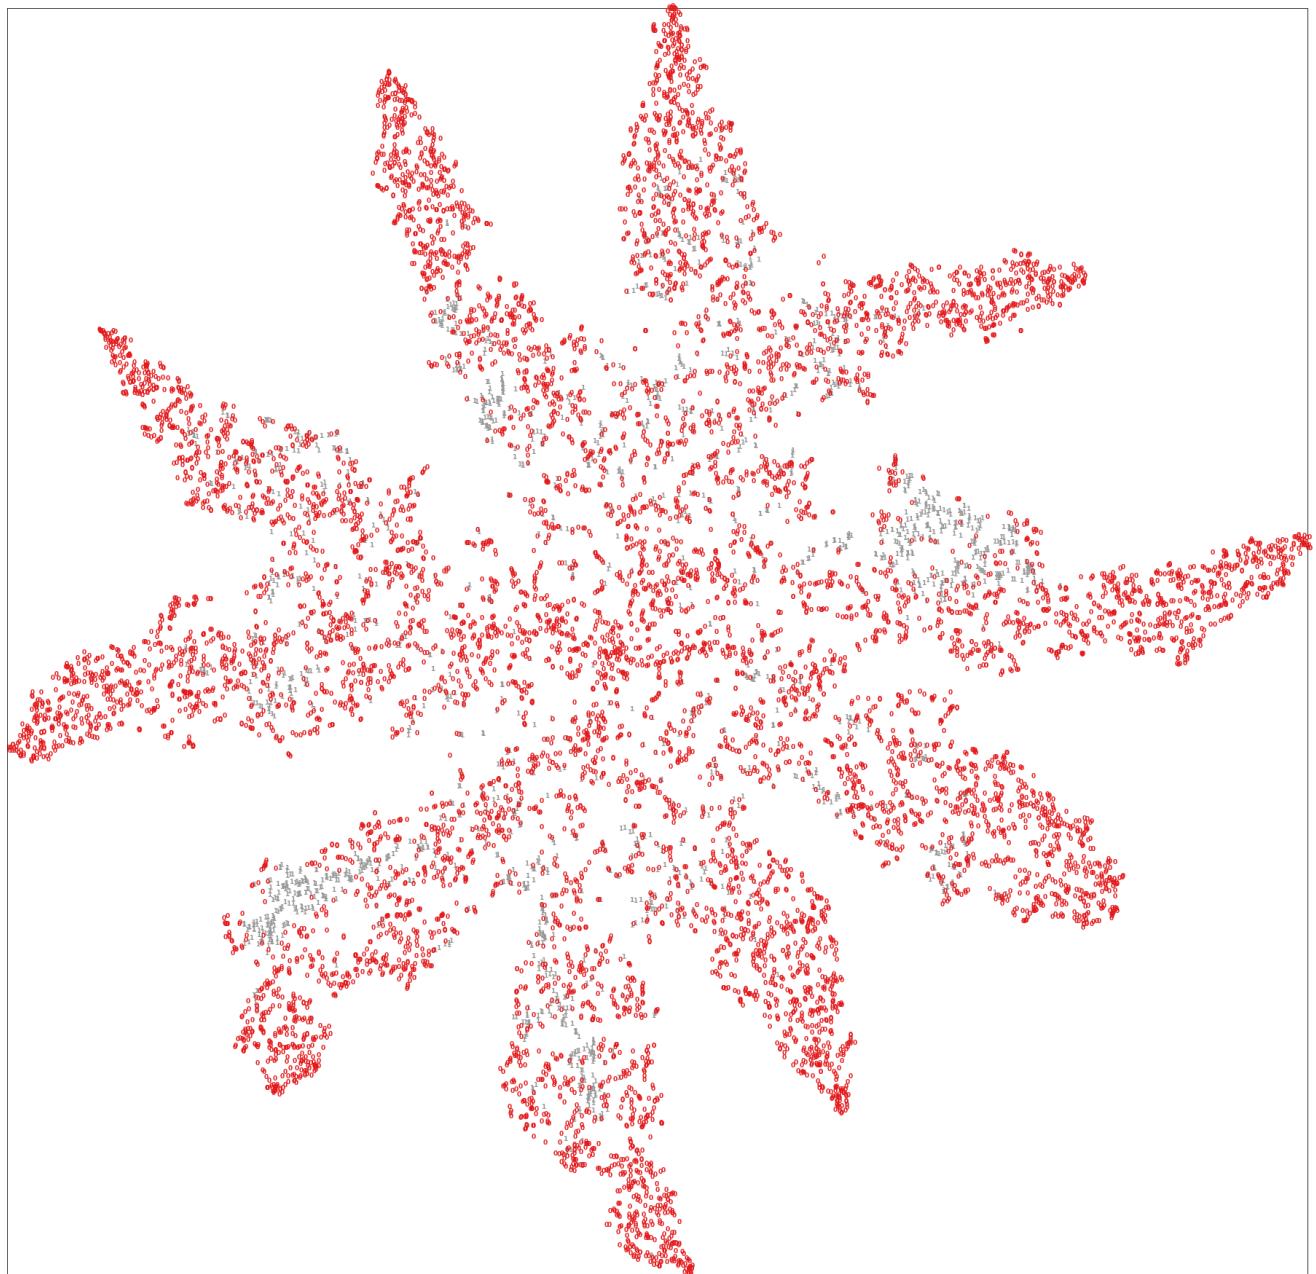
(b)



- mnistm->usps
 - (a)



(b)



5.

On usps->svhn, if we use models that are too deep, the performance would be bad. I grayscaled all images if source or domain is usps. I used a model that has featureExtractor, domainClassifier and labelPredictor, and a GRL layer. The model is simple but enough to pass all baselines.

Bonus

1.

	MNSITM → USPS	SVHN → MNIST-M	USPS → SVHN
DANN	0.7643	0.4918	0.3041
improve	0.9147	0.6805	0.4810

2.

I used auto-augment and I implemented a MCD model from this https://github.com/mil-tokyo/MCD_DA (https://github.com/mil-tokyo/MCD_DA).

I guess the reason that the new one improves because the model and loss function are better and autoaugment let the distribution of source and target more likely to be aligned.